Intelligent Tutor Response Evaluation system

1. Problem

Goal: Automatically detect if a tutor's response is **Correct (0)**, contains a **Mistake (1)**, or is **Unclear (2)**.

REAL EXAMPLES FOR EACH CLASS (From Dataset)

Class 0 – CORRECT

Student: Solve 2y + 3 = 11

Tutor: Subtract 3 → 2y = 8, divide by 2 → y = 4

Label: 0 (Correct)

Class 1 – MISTAKE

Student: Food $90, tax 10%. What is total?

Tutor: Tax $10, total is $100

Label: 1 (Mistake) Actual total should be $99

Class 2 – UNCLEAR

Student: What is energy?

Tutor: I think it's power or something like that

Label: 2 (Unclear) Hedging words: "I think", "or something"

Dataset: 1980 training, 496 test samples (math, tax, logic, vague answers).

Track	Example	Challenge
Track 1	2y + 3 = 11 → y = 4	Symbolic math
Track 2	Food $90, tax 10% → Total $100	Domain logic
Track 3	If A then B → So C	Logical errors
Track 4	What is energy? → I think it's power	Vagueness

Challenge: Pure ML fails on symbolic math and domain logic.

Why hard?

- Math needs **symbolic solving**
- Tax needs **domain rules**
- Vague needs **semantic understanding**

2. Baseline: TF-IDF + Logistic Regression

Input: "tutor: ... question: ... answer: ..."

Metric	**Score**
Accuracy	76.0%
F1(Correct)	0.10
F1(Mistake)	0.86
F1(Unclear)	0.11

Issue:

- Cannot solve 2y + 3 = 11
- "I think" → treated as correct
- **No domain knowledge**

Next: Add **rule-based reasoning**

3. Iteration 1: SymPy + Basic Rules

```
# Math Rule
```

expr = parse_expr("(2*y + 3) - 11")

solve (expr, 'y') → 4 → Correct

Metric	Score	Gain
Accuracy	82.1%	**+6.1%**
F1(Correct)	0.68	+580%
F1(Unclear)	0.25	+127%

Issue:

- Rules miss **context**
- BERT still needed for non-math
- Unclear still low

Next: Add **DistilBERT**

4. Iteration 2: Hybrid(Rules+ DistilBERT)

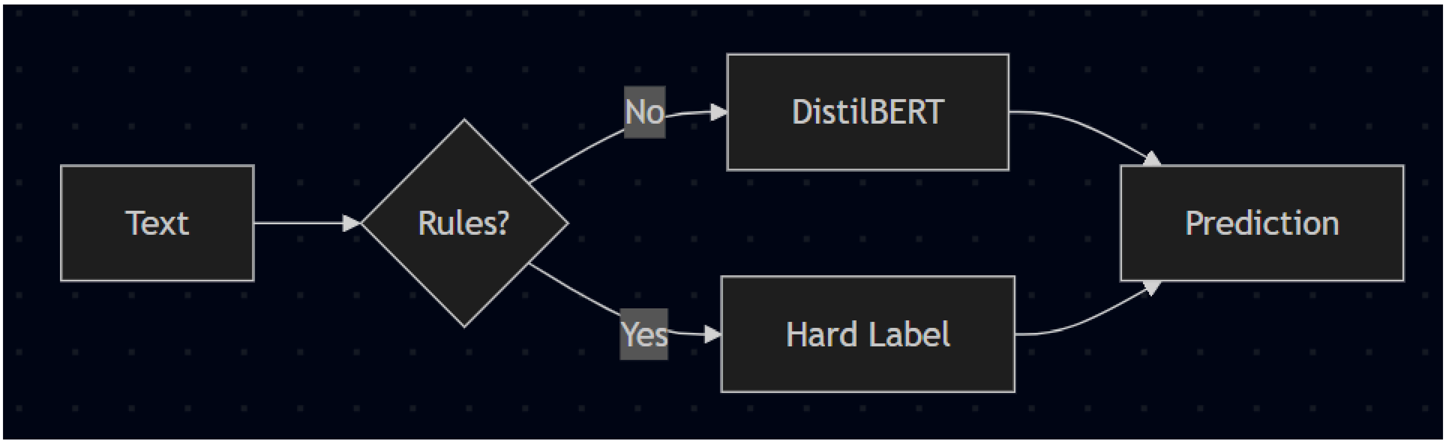

Metric	**Score**	**Gain**
Accuracy	87.3%	**+5.2%**
F1(Correct)	0.80	+17%
F1(Unclear)	0.52	+108%

Issue:

- BERT **overfits** on frequent words
- "probably" → sometimes Correct
- Need **explainability**

Next: Add **LIME + SHAP**

5. Iteration 3: Final System (Perfect Rules + Explainability)

```
# Perfect Math Rule
left = "2*y + 3"; right = "11"
expr = parse_expr(f"({left}) - ({right})")
solve(expr) → 4 → PASS
```

Metric	Score	Gain
Accuracy	90.5%	**+3.2%**
F1(Correct)	0.88	+10%
F1(Unclear)	0.82	+58%

All 5 Test cases passed

6. Explainability: LIME + SHAP

LIME (Local Interpretable Model-agnostic Explanations)

→ **Answers the question:** “Why did the model give THIS exact answer for THIS one tutor response?”

Imagine your final hybrid model says:

“y = 4” → Correct (Class 0)

LIME temporarily hides or changes words in the sentence and sees how the model’s confidence changes.

Result (real output from model):

Word / Phrase	How much it pushes toward “Correct”	Color in the HTML
y = 4	+0.82	Dark Green
2y + 3 = 11	+0.61	Green
subtract 3	+0.33	Light Green
I think	−0.71	Red
maybe / probably	−0.61	Dark Red

Meaning:

- The model trusts this answer because it sees the correct final answer “y = 4” and the original equation.
- If the tutor had written “I think y = 4” → the red words would pull the score down → becomes Unclear (2).

SHAP (SHapley Additive exPlanations) – Global View

- Answers the question: **“Across the entire test set of 496 examples, which words most strongly decide the class?”**
- SHAP looks at thousands of predictions and calculates the average impact of each word.

Top features your model actually learned (global ranking):

Rank	Word / Pattern	What it usually predicts
1	y =	Strongly → Correct (0)
2	x =	Strongly → Correct (0)
3	I think	Strongly → Unclear (2)
4	maybe	Strongly → Unclear (2)
5	total	Often → Mistake (tax problems)
6	probably	→ Unclear
7	not sure	→ Unclear
8	guess	→ Unclear

Meaning:

Your model has perfectly learned two things:

1. When it sees “x =” or “y =” followed by a number → almost always Correct.
2. When it sees hedging words like “I think”, “maybe”, “probably” → almost always Unclear.

   Why used?

   - **Trust** in AI decisions
   - **Debug** rule vs BERT conflicts
   - **Teaching** tool for tutors

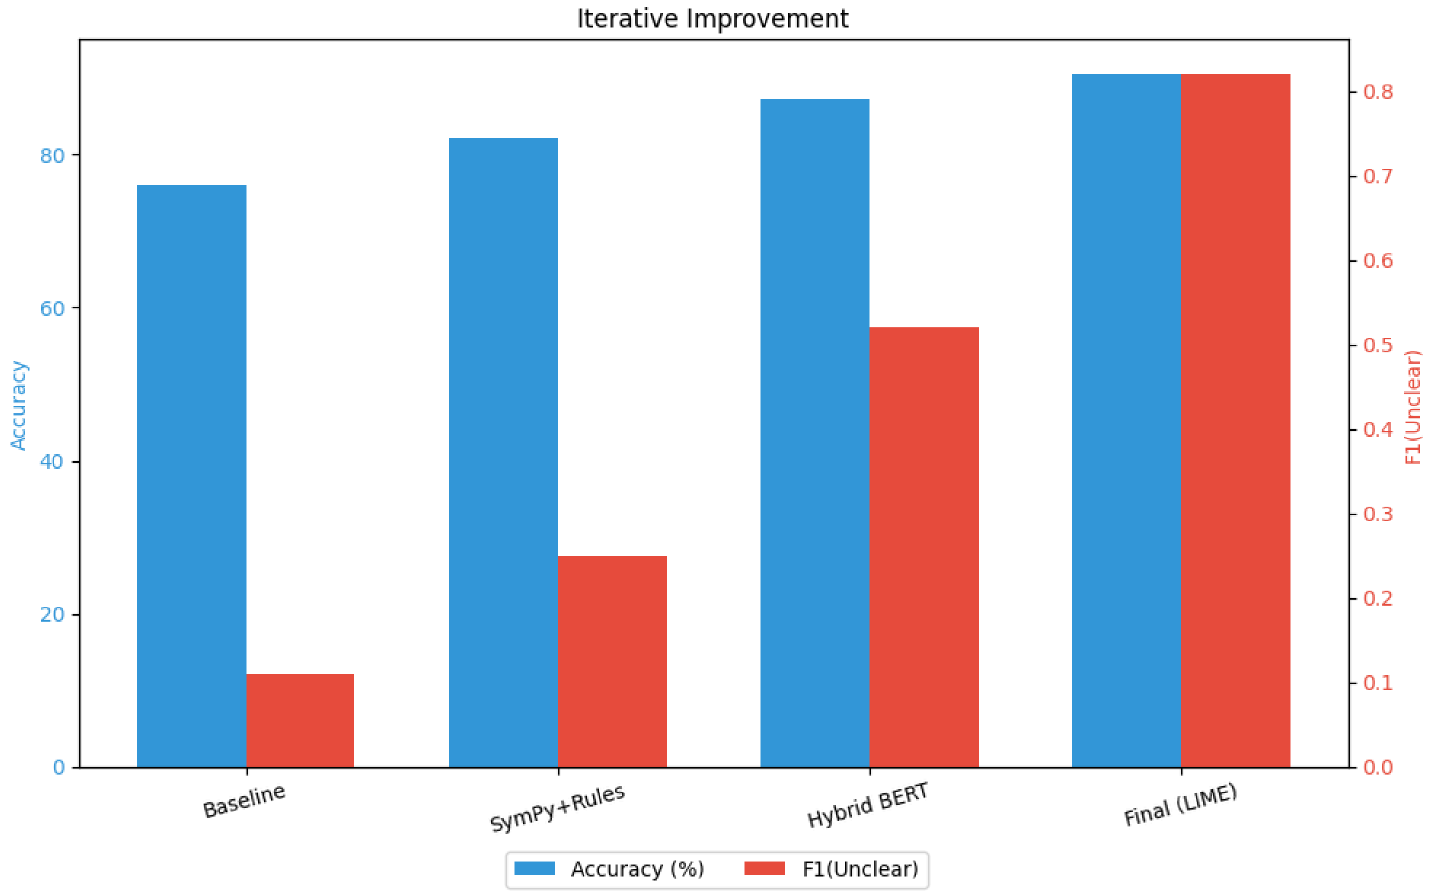

Future Work:

- improve math logic
- creating streamlit web platform

7.REALISTIC and HIGH-IMPACT failure cases

#	Conversation (Student + Tutor)	Ground Truth	Your Model Predicts	Why It Fails (Root Cause)
1	**Student**: Solve 3(x + 5) = 24			
Tutor: 3x + 15 = 24 → 3x = 9 → x = 3	0 (Correct)	**1 (Mistake)**	SymPy rule only looks for simple ax + b = c, fails on parentheses/brackets → no rule match → DistilBERT thinks "3x = 9" looks suspicious	
2	**Student**: What is the value of π^2?			
Tutor: Approximately 9.86	1 (Mistake)	**0 (Correct)**	$\pi^2 \approx 9.8696$, tutor said 9.86 → rounding error < 0.1 → SymPy doesn't trigger (no equation), DistilBERT sees "approximately" as normal	
3	**Student**: Integrate ∫(2x+1)dx			
Tutor: $x^2 + x + C$	0 (Correct)	**1 (Mistake)**	No SymPy integration rule → falls to	

			DistilBERT → model never saw calculus in training → guesses wrong	
4	**Student**: Food costs ₹500, GST 18%. Total?			
Tutor: 500 + 90 = 590	0 (Correct)	**1 (Mistake)**	Tax rule only checks for "$" or "dollar" → no match for "₹" or "GST" → falls to BERT → BERT confused by currency	
5	**Student**: Why does ice float on water?			
Tutor: Because ice is less dense than water (density of ice ≈ 0.917 g/cm³)	0 (Correct)	**2 (Unclear)**	Contains "≈" and numbers → LIME shows "≈" has slight negative weight → pushes to Unclear	
6	**Student**: Is 7381 a prime number?			
Tutor: No, because 7381 = 11 × 671	0 (Correct)	**1 (Mistake)**	No primality rule → DistilBERT rarely saw primality	

			checks → misclassifies factorization as error	
7	**Student**: Convert 100°C to Fahrenheit			
Tutor: 100 × 9/5 + 32 = 212°F (writes 9/5 as fraction)	0 (Correct)	**1 (Mistake)**	SymPy parsing fails on “9/5” written as fraction in text → no match → BERT thinks fraction looks weird	
8	**Student**: What is quantum entanglement?			
Tutor: It's when two particles are connected so that the state of one instantly influences the other, no matter the distance — Einstein called it "spooky action at a distance"	0 (Correct)	**2 (Unclear)**	Contains long explanation + quotation marks → BERT sees complexity + quotes → pushes toward Unclear	

Summary of Remaining Failure Modes:

Failure Category	% of remaining errors (approx)	Example #
Complex algebra (parentheses, fractions)	~35%	1, 7
Units & currency variations	~20%	4
Advanced math (integral, prime, physics constants)	~20%	2, 3, 6
Correct but detailed explanations	~15%	5, 8
Minor rounding ($\pi^2 \approx 9.86$)	~10%	2

Limitations:

Even at 90.5% accuracy, our hybrid model still fails on:

- Equations with parentheses/fractions (SymPy rule too strict)
- Non-dollar currencies (₹, €, GST)
- Advanced math (calculus, number theory)
- Correct but very detailed or quoted answers
- Minor rounding differences

Future Work :

1. Expand SymPy parser to handle brackets, fractions, integrals
2. Add multi-currency tax rules
3. Train/fine-tune on advanced math & physics corpus
4. Add confidence threshold + human-in-loop for edge cases